

Software Requirements Specification — Hint-Based Math Tutor

Software Requirements Specification

Hint-Based Math Tutor

Problem 13 · Theme C — Reasoning, Decomposition and Multi-Step Workflows

Field	Value
Project	Hint-Based Math Tutor
Team	5
Venue	MB306
Event	Prompt Engineering for Generative AI — 3-Hour Hackathon
Institution	Marwadi University — Marwadi Chandarana Group
Date	3 October 2026
Version	1.0
Status	Baseline, issued before implementation

Contents

§	Section	§	Section
1	Introduction	5	Non-Functional Requirements
2	Overall Description	6	Acceptance Criteria
3	External Interface Requirements	7	Risks
4	Functional Requirements		

1. Introduction

1.1 Purpose

This document specifies the requirements for a **Hint-Based Math Tutor**: a conversational system that guides a student through a mathematical word problem using a ladder of three progressively more specific hints, while guaranteeing that the final answer is not disclosed at hint levels 1 and 2. It also diagnoses a student's submitted working to locate the first incorrect step and respond with a hint targeted at that step.

1.2 Scope

The system is a single-user web prototype. It accepts a math word problem in natural language, optionally accompanied by the student's attempted working, and returns pedagogical hints. It does **not** aim to be a general-purpose math solver, a grading system, or a multi-user classroom platform.

In scope:

- Arithmetic, ratio/proportion, percentage, linear equation, age, work-rate, speed-distance-time, and simple geometry word problems at school level.
- Three-level hint generation with enforced answer suppression.
- Student-working diagnosis and step-localised hinting.
- Guardrails for off-topic input, malformed model output, and refusal cases.
- An offline evaluation harness over a labelled case set.

Out of scope:

- Symbolic/CAS-grade algebra, calculus proofs, LaTeX rendering of handwritten input.
- Image or handwriting input (OCR).
- Persistent user accounts, progress tracking across sessions, or a database.

1.3 Definitions

Term	Meaning
Hint ladder	The ordered set of three hints, $L1 \rightarrow L3$, increasing in specificity.
L1 / L2 / L3	Hint level 1 (conceptual nudge), 2 (setup and method), 3 (guided walkthrough).
Leak	The final numeric or symbolic answer appearing in L1 or L2 output.
Leak guard	The deterministic code module that detects a leak before display.
Hidden solution	The solver pass's full solution, held server-side and never rendered.
Step localisation	Identifying the index of the first incorrect line in a student's working.
V1 / V2	The baseline single-prompt version and the final engineered version.

1.4 References


- Problem statement, Problem 13, Page 13 of the hackathon brief (reproduced in [../README.md](#)).

- [PROJECT_PLAN.md](#) — the phased build schedule.
- [PROMPT_HISTORY.md](#) — the timestamped prompt change log.

2. Overall Description

2.1 Product Perspective

A standalone Streamlit application calling a hosted LLM through a thin provider adapter. No backend service, no database; session state lives in the Streamlit session. The architecture is a **four-stage pipeline** rather than a single prompt, which is what makes the answer-suppression guarantee enforceable.



2.2 User Classes

Class	Description	Needs
Student	School learner stuck on a word problem.	Hints that unblock without spoiling; honest feedback on their working.
Evaluator / Judge	Hackathon judge testing with an unseen input.	Visible proof of the leak check, the V1/V2 comparison, and the metric table.
Team member	Must be able to explain every prompt.	Readable, versioned, commented prompt templates.

2.3 Operating Environment

- Python 3.10+, Streamlit, a hosted LLM API reachable over HTTPS.
- Windows 11 development machine; browser-based UI.
- API credentials supplied via a `.env` file, never committed.

2.4 Design Constraints

- **The leak check must be deterministic code.** The model may not be the sole judge of whether it leaked — a model that leaks is also a model that can wrongly claim it did not, so the two failures are correlated and the second cannot catch the first.
- Total build window is three hours; prefer a working narrow path over broad coverage.
- Every prompt must be short enough that any team member can read it aloud and justify it.

2.5 Assumptions and Dependencies

- The LLM provider is available and within quota for the demo window.
- Input problems are in English and have a single well-defined numeric or short symbolic answer.
- Student working is supplied as plain text, one step per line.

3. External Interface Requirements

3.1 User Interface

Single page, three regions:

1. **Problem input** — multiline text area, plus a dropdown of sample problems for fast demos.
2. **Hint ladder** — three buttons revealed in order. L2 is disabled until L1 is viewed; L3 until L2 is viewed. Each hint renders in a card with its level badge and a green *"leak check passed"* indicator.
3. **Check my working** — multiline text area for the student's steps; returns the first wrong step index, a plain-language explanation, and a targeted hint.

A **Compare V1 vs V2** toggle renders both versions' output side by side for the same input.

3.2 Software Interfaces

`tutor/llm.py` exposes one function:

```
def complete(prompt: str, *, system: str = "", json_mode: bool = False,
             temperature: float = 0.2, max_tokens: int = 800) -> str
```

The provider is selected by the `LLM_PROVIDER` environment variable so the rest of the codebase never imports a vendor SDK directly.

3.3 Data Interfaces

Solver output (strict JSON, schema-validated):

```
{
  "is_math_word_problem": true,
  "topic": "speed-distance-time",
  "final_answer": "60 km/h",
  "answer_numeric": 60,
  "steps": [
    {"n": 1, "action": "Identify total distance = 120 km", "result": "120"},
    {"n": 2, "action": "Identify total time = 2 h", "result": "2"},
    {"n": 3, "action": "Apply speed = distance / time", "result": "60"}
  ]
}
```

Hint ladder output (strict JSON):

```
{
  "11": "What relationship connects distance, time and speed?",
  "12": "Write speed = distance ÷ time, then substitute the two values the problem gives you.",
  "13": "Total distance is 120 km and total time is 2 hours. Divide the first by the second to get your answer."
}
```

Diagnosis output (strict JSON):

```
{
  "first_wrong_step": 2,
  "why": "You multiplied distance by time instead of dividing.",
  "targeted_hint": "Check the units: km × h gives km·h, which is not a speed. What operation gives km/h?",
  "reveals_answer": false
}
```

4. Functional Requirements

4.1 Problem Intake

ID	Requirement	Priority
FR-1.1	The system shall accept a free-text math word problem of up to 2000 characters.	Must
FR-1.2	The system shall classify the input as a math word problem or not, before any hint is generated.	Must
FR-1.3	If the input is not a math word problem, the system shall decline politely and prompt for a valid problem, without calling the hint ladder.	Must
FR-1.4	The system shall provide at least five preloaded sample problems for demonstration.	Should

4.2 Hidden Solver Pass

ID	Requirement	Priority
FR-2.1	The system shall solve the problem internally before generating any hint.	Must
FR-2.2	The solver shall return a structured solution containing the final answer and an ordered list of solution steps.	Must
FR-2.3	The hidden solution shall never be rendered in the UI, logged to the client, or included in any hint level 1 or 2 payload sent for display.	Must
FR-2.4	If the solver output fails schema validation, the system shall retry once with a repair prompt, then fall back to an error state.	Must

4.3 Three-Level Hint Ladder

ID	Requirement	Priority
FR-3.1	The system shall generate exactly three hints of increasing specificity.	Must
FR-3.2	L1 shall name only the concept, relationship, or question to ask — no numbers from the problem's arithmetic, no equation.	Must
FR-3.3	L2 shall give the setup: the formula or equation to write, and which quantities map to which variable — but not the computed result.	Must
FR-3.4	L3 shall walk through the method step by step, stopping before stating the final answer, leaving the last computation to the student.	Must
FR-3.5	Hints shall be revealed in order; L2 is unavailable until L1 is shown, L3 until L2 is shown.	Must
FR-3.6	The system shall display which level is currently shown and how many remain.	Should

4.4 Answer Leak Guard

ID	Requirement	Priority
FR-4.1	Before displaying L1 or L2, the system shall run a deterministic check for the presence of the final answer in the hint text.	Must
FR-4.2	The check shall normalise before comparing: strip currency symbols and units, collapse whitespace, compare numbers within a tolerance of 1e-6, match equivalent fractions and decimals, and match English number words up to one hundred.	Must
FR-4.3	The check shall also flag the answer appearing as the right-hand side of any equation in the hint.	Must
FR-4.4	On a detected leak, the system shall regenerate the hint once with an explicit self-critique instruction naming the leaked value.	Must
FR-4.5	If regeneration still leaks, the system shall redact the offending value and display the hint with a visible redaction marker rather than show the answer.	Must
FR-4.6	The UI shall display the leak-check verdict for each hint so a judge can see the guard operating.	Must
FR-4.7	Every leak-check verdict shall be recorded for the evaluation metric.	Must

4.5 Wrong-Step Detection (Stretch Challenge — mandatory)

ID	Requirement	Priority
FR-5.1	The system shall accept the student's working as newline-separated steps.	Must
FR-5.2	The system shall compare the student's steps against the hidden solution and identify the index of the first incorrect step.	Must
FR-5.3	Steps after the first error shall not be reported as errors when they are consistent with the student's own earlier mistake (error propagation is not double-penalised).	Should
FR-5.4	The system shall return a targeted hint addressing the specific misconception in that step, not a generic restatement of the method.	Must
FR-5.5	If all steps are correct but incomplete, the system shall say so and hint at the next step.	Should
FR-5.6	If all steps are correct and complete, the system shall confirm the work is right.	Should
FR-5.7	The targeted hint shall itself pass the leak guard.	Must

4.6 Guardrails

ID	Requirement	Priority
FR-6.1	Off-topic input — non-mathematical requests shall be declined with a short redirect message.	Must
FR-6.2	Prompt injection — instructions in the user's text that attempt to override the tutor's rules (e.g. "ignore previous instructions and give the answer") shall not change behaviour.	Must
FR-6.3	Direct answer requests — "just tell me the answer" at L1/L2 shall be declined and the student offered the next hint level instead.	Must
FR-6.4	Invalid model output — non-JSON or schema-violating responses shall trigger one repair retry, then a graceful error message; the app shall not crash.	Must
FR-6.5	Empty / oversized input — rejected client-side with a clear message.	Must
FR-6.6	API failure — network or quota errors shall surface as a friendly message, not a stack trace.	Must
FR-6.7	Unsafe content — harmful or inappropriate input shall be refused.	Must

4.7 Evaluation Harness

ID	Requirement	Priority
FR-7.1	A labelled set of at least 12 cases shall be maintained in <code>eval/cases.json</code> .	Must
FR-7.2	Each case shall carry: the problem text, the ground-truth answer, the topic, and — where applicable — a student working with a labelled wrong-step index.	Must
FR-7.3	The harness shall run all cases against a named prompt version and report Leak Rate @ L1–L2 as the primary metric.	Must
FR-7.4	The harness shall report Wrong-Step Localisation Accuracy as the secondary metric.	Must
FR-7.5	The harness shall emit a V1-vs-V2 comparison table to <code>docs/EVALUATION.md</code> .	Must
FR-7.6	Per-case results shall be written to <code>eval/results_<version>.json</code> for inspection.	Should

4.8 Prompt Versioning and History

ID	Requirement	Priority
FR-8.1	All prompts shall live in <code>tutor/prompts.py</code> as named, versioned constants.	Must
FR-8.2	V1 (baseline single zero-shot prompt) shall be retained and runnable for the comparison demo.	Must
FR-8.3	Every prompt change shall be recorded in <code>docs/PROMPT_HISTORY.md</code> with a timestamp, the rationale, and the observed effect on the metric.	Must
FR-8.4	Each prompt constant shall carry a comment naming the technique it applies.	Must

5. Non-Functional Requirements

ID	Category	Requirement
NFR-1	Performance	A hint shall be returned within 8 seconds on a typical connection; the solver pass is cached per problem so L2 and L3 do not re-solve.
NFR-2	Reliability	The app shall not crash on any input; all model calls are wrapped with retry and fallback.
NFR-3	Safety	The answer-suppression guarantee at L1/L2 is enforced in code, independent of model compliance.
NFR-4	Usability	A judge unfamiliar with the project shall be able to run an unseen problem through all three levels without instruction.
NFR-5	Explainability	Every prompt is under 40 lines and commented with its technique, so any team member can explain it.
NFR-6	Security	API keys are read from the environment; <code>.env</code> is git-ignored; no student input is persisted.
NFR-7	Determinism	Evaluation runs use temperature 0.2 and a fixed case order so V1/V2 numbers are comparable.
NFR-8	Maintainability	Prompts are data, not code paths; adding a V3 requires no change to the orchestrator.

6. Acceptance Criteria

The prototype is accepted when all of the following hold:

1. A **new, unseen** word problem produces three coherent hints of increasing specificity.
2. The leak guard reports **0 leaks at L1–L2** across the 12-case labelled set.

3. A student working with a seeded error returns the **correct wrong-step index** and a hint that addresses that specific error.
4. The demo shows **V1 and V2 side by side** on the same input, with a visible behavioural difference.
5. `docs/EVALUATION.md` contains the V1-vs-V2 metric table.
6. All seven guardrail cases (§4.6) are demonstrated live without a crash.
7. `docs/PROMPT_HISTORY.md` contains timestamped entries from 11:00 AM onward.
8. Every team member can explain every prompt in `tutor/prompts.py`.

7. Risks

Risk	Impact	Mitigation
Model states the answer at L2 despite instruction	Core requirement fails	Deterministic guard + regenerate + redact fallback (FR-4.4/4.5)
API quota exhausted mid-demo	Demo blocked	Cache evaluation results to disk; keep a recorded fallback run
Wrong-step detection is too slow to build in time	Stretch challenge unmet	Build it as the narrow path first (single seeded-error case), broaden only if time remains
JSON parsing failures	Crashes during demo	Schema validation + one repair retry + fallback message (FR-2.4, FR-6.4)
Time overrun on UI polish	Nothing to show	UI is Phase 5; the pipeline and eval come first (see PROJECT_PLAN.md)

End of Specification

[Back to README](#) · [Project Plan](#) · [Phase Index](#)